

EEE-6512: Image Processing and Computer Vision

November 13, 2018

Lecture #12: Model Fitting

Damon L. Woodard, Ph.D.

Dept. of Electrical and Computer
Engineering

dwoodard@ece.ufl.edu

Fitting

- We've learned how to detect edges, and corners. Now what?
- We would like to form a higher-level, more compact representation of the features in the image by grouping multiple features according to a simple model




Fitting: Main idea

- **Choose a parametric model to represent a set of features**
- **Membership criterion is not local**
 - Can't tell whether a point belongs to a given model just by looking at that point
- **Three main questions:**
 - What model represents this set of features best?
 - Which of several model instances gets which feature?
 - How many model instances are there?
- **Computational complexity is important**
 - It is infeasible to examine every possible set of parameters and every possible combination of features

Fitting


Choose a *parametric model* to represent a set of features



simple model: lines



simple model: circles



complicated model: car




complicated model: face shape

Fitting -Design challenges

- Design a suitable **goodness of fit** measure
 - Similarity should reflect application goals
 - Encode robustness to outliers and noise
- Design an **optimization** method
 - Avoid local optima
 - Find best parameters quickly



Example: Line fitting

- Why fit lines?
Many objects characterized by presence of straight lines



- Wait, why aren't we done just by running edge detection?

Difficulty of line fitting



- **Extra** edge points (clutter), multiple models:
 - which points go with which line, if any?
- Only some parts of each line detected, and some parts are **missing**:
 - how to find a line that bridges missing evidence?
- **Noise** in measured edge points, orientations:
 - how to detect true underlying parameters?

Fitting: Overview (cont.)

- **If we know which points belong to the line, how do we find the “optimal” line parameters?**

Least squares

- **What if there are outliers?**

RANSAC

- **What if there are many lines?**

Voting methods: RANSAC, Hough transform

Line Fitting

Fitting: Methods


- Global optimization / Search for parameters
 - Ordinary Least Squares
 - Total Least Squares
- Hypothesize and test
 - Hough transform
 - RANSAC

Ordinary Least Squares / Total Least
Squares

Least squares line fitting

- Data: $(x_1, y_1), \dots, (x_n, y_n)$
- Line equation: $y_i = mx_i + b$
- Find (m, b) to minimize

$$E = \sum_{i=1}^n (y_i - mx_i - b)^2$$



$$E = \sum_{i=1}^n \left(\begin{bmatrix} x_i & 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} - y_i \right)^2 = \left\| \begin{bmatrix} x_1 \\ \vdots \\ x_n \\ 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} - \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix} \right\|^2 = \|\mathbf{A}\mathbf{p} - \mathbf{y}\|^2$$

$$= \mathbf{y}^T \mathbf{y} - 2(\mathbf{A}\mathbf{p})^T \mathbf{y} + (\mathbf{A}\mathbf{p})^T (\mathbf{A}\mathbf{p})$$


$$\frac{dE}{dp} = 2\mathbf{A}^T \mathbf{A}\mathbf{p} - 2\mathbf{A}^T \mathbf{y} = 0$$

Matlab: `p = A \ y;`

$$\mathbf{A}^T \mathbf{A}\mathbf{p} = \mathbf{A}^T \mathbf{y} \Rightarrow \mathbf{p} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{y}$$


Problem with “vertical” least squares

- Not rotation-invariant
- Fails completely for vertical lines



Total least squares

- Distance between point (x_n, y_n) and line $ax+by=d$ ($a^2+b^2=1$): $|ax + by - d|$
- Find (a, b, d) to minimize the sum of squared perpendicular distances




$$E = \sum_{i=1}^n (ax_i + by_i - d)^2$$

Total least squares

• Distance between point (x_n, y_n) and line $ax+by=d$ ($a^2+b^2=1$): $|ax + by - d|$

• Find (a, b, d) to minimize the sum of squared perpendicular distances

$$E = \sum_{i=1}^n (ax_i + by_i - d)^2$$



$$\frac{\partial E}{\partial d} = \sum_{i=1}^n -2(ax_i + by_i - d) = 0$$

$$d = \frac{a}{n} \sum_{i=1}^n x_i + \frac{b}{n} \sum_{i=1}^n y_i = a\bar{x} + b\bar{y}$$


$$E = \sum_{i=1}^n (a(x_i - \bar{x}) + b(y_i - \bar{y}))^2 = \left\| \begin{bmatrix} x_1 - \bar{x} & y_1 - \bar{y} \\ \vdots & \vdots \\ x_n - \bar{x} & y_n - \bar{y} \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix} \right\|^2 = (UN)^T (UN)$$

$$\frac{dE}{dN} = 2(U^T U)N = 0$$

Solution to $(U^T U)N = 0$, subject to $\|N\|^2 = 1$: eigenvector of $U^T U$ associated with the smallest eigenvalue (least squares solution to *homogeneous linear system* $UN = 0$)


Least squares: Robustness to noise

Least squares fit to the red points:



Least squares: Robustness to noise

Least squares fit with an outlier:



Problem: squared error heavily penalizes outliers

Least squares (global) optimization

Good

- Clearly specified objective
- Optimization is easy

Bad

- Sensitive to outliers
 - Bad matches, extra points
- Doesn't allow you to get multiple good fits
 - Detecting multiple objects, lines, etc.

Other ways to search for parameters (for when no closed form solution exists)

- **Line search**

1. For each parameter, step through values and choose value that gives best fit
2. Repeat (1) until no parameter changes

- **Grid search**

1. Propose several sets of parameters, evenly sampled in the joint set
2. Choose best (or top few) and sample joint parameters around the current best; repeat

- **Gradient descent**

1. Provide initial position (e.g., random)
2. Locally search for better parameters by following gradient

Hypothesize and Test/Hough Transform

Hypothesize and Test

1. Propose parameters

- Try all possible
- Each point votes for all consistent parameters
- Repeatedly sample enough points to solve for parameters

2. Score the given parameters

Number of consistent points, possibly weighted by distance

3. Choose from among the set of parameters

Global or local maximum of scores

4. Possibly refine parameters using inliers

Hough Transform: Outline


1. Create a grid of parameter values
2. Each point votes for a set of parameters, incrementing those values in grid
3. Find maximum or local maxima in grid

Connection Between Image and Hough Spaces



Hough transform





Given a set of points, find the curve or line that explains the data points best



$$y = m x + b$$




Hough transform



Polar representation for lines

Issues with usual (m,b) parameter space: can take on infinite values, undefined for vertical lines.



d : perpendicular distance
from line to origin

θ : angle the perpendicular
makes with the x-axis

$$x \cos \theta - y \sin \theta = d$$

Point in image space \rightarrow sinusoid segment in Hough space

Image Parameter Spaces

Image Space	Parameter Space
Lines	Points
Points	Lines/Sinusoid
Collinear Points	Intersecting Lines

Hough transform algorithm

Using the polar parameterization:


$$x \cos \theta - y \sin \theta = d$$

Basic Hough transform algorithm


1. Initialize $H[d, \theta] = 0$
2. for each edge point $I[x, y]$ in the image
 for $\theta = [\theta_{\min} \text{ to } \theta_{\max}]$ // some quantization
 $d = x \cos \theta - y \sin \theta$
 $H[d, \theta] += 1$
3. Find the value(s) of (d, θ) where $H[d, \theta]$ is maximum
4. The detected line in the image is given by

$$d = x \cos \theta - y \sin \theta$$

H: accumulator array (votes)





Hough transform - experiments



Need to adjust grid size or smooth



Hough transform - experiments



Issue: spurious peaks due to uniform noise

Dealing with noise

- **Choose a good grid / discretization**
 - Too coarse: large votes obtained when too many different lines correspond to a single bucket
 - Too fine: miss lines because some points that are not exactly collinear cast votes for different buckets
- **Increment neighboring bins (smoothing in accumulator array)**
- **Try to get rid of irrelevant features**
 - Take only edge points with significant gradient magnitude

Hough Algorithm Extensions

Extensions

Extension 1: Use the image gradient

1. same
2. for each edge point $I[x,y]$ in the image


$$\theta = \text{gradient at } (x,y)$$

$$d = x \cos \theta - y \sin \theta$$

$$H[d, \theta] += 1$$

3. same
4. same

(Reduces degrees of freedom)



$$\theta = \tan^{-1} \left(\frac{\frac{\partial f}{\partial y}}{\frac{\partial f}{\partial x}} \right)$$

Extension 2

- give more votes for stronger edges

Extension 3

- change the sampling of (d, θ) to give more/less resolution

Extension 4

- The same procedure can be used with circles, squares, or any other shape


Hough Transform for Circles

Hough transform for circles

- Circle: center (a,b) and radius r

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For a fixed radius r , unknown gradient direction




Hough transform for circles

- Circle: center (a,b) and radius r

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For a fixed radius r , unknown gradient direction




Hough transform for circles

- Circle: center (a,b) and radius r

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For an unknown radius r , unknown gradient direction




Hough transform for circles

- Circle: center (a,b) and radius r

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For an unknown radius r , unknown gradient direction



Hough transform for circles

- Circle: center (a,b) and radius r

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For an unknown radius r , **known** gradient direction





Image space



Hough space

Hough transform for circles

For every edge pixel (x,y) :

For each possible radius value r :

For each possible gradient direction θ :

$a = x - r \cos(\theta)$ // column

$b = y - r \sin(\theta)$ // row

$H[a,b,r] += 1$

end

end

Optimization of Circle Hough Transform

Use Edge Direction (Eliminates Radius)


$$b = a * \tan(\theta) - x * \tan(\theta) + y$$

Example: detecting circles with Hough


Original



Edges




Votes: Penny




Note: a different Hough transform (with separate accumulators) was used for each circle radius (quarters vs. penny).

Example: detecting circles with Hough


Original



Edges



Votes: Quarter




Combined detections

Example: iris detection



Gradient+threshold




Hough space
(fixed radius)



Max detections

Voting: practical tips

- Minimize irrelevant tokens first
- Choose a good grid / discretization



- Vote for neighbors, also (smoothing in accumulator array)
- Use direction of edge to reduce parameters by 1
- To read back which points voted for “winning” peaks, keep tags on the votes.

Parameters for analytic curves

Analytic Form	Parameters	Equation
Line	ρ, θ	$x\cos\theta + y\sin\theta = \rho$
Circle	x_0, y_0, ρ	$(x-x_0)^2 + (y-y_0)^2 = r^2$
Parabola	x_0, y_0, ρ, θ	$(y-y_0)^2 = 4\rho(x-x_0)$
Ellipse	x_0, y_0, a, b, θ	$(x-x_0)^2/a^2 + (y-y_0)^2/b^2 = 1$


Speed of Hough Transform

The computation time grows exponentially as the number of parameters increases

1. Image → Canny




2. Canny \rightarrow Hough votes



3. Hough votes \rightarrow Edges

Find peaks and post-process



Hough transform: pros and cons

Pros

- All points are processed independently, so can cope with occlusion, gaps
- Some robustness to noise: noise points unlikely to contribute *consistently* to any single bin
- Can detect multiple instances of a model in a single pass

Cons

- Complexity of search time increases exponentially with the number of model parameters
- Non-target shapes can produce spurious peaks in parameter space
- Quantization: can be tricky to pick a good grid size

Addressing Noise: RANSAC

RANSAC

- **RAN**dom **SA**mples **C**onsensus
- Approach: we want to avoid the impact of outliers, so let's look for "inliers", and use only those.
- Intuition: if an outlier is chosen to compute the current fit, then the resulting line won't have much support from rest of the points.

RANSAC


RANSAC loop:

1. Randomly select a *seed group* of points on which to base transformation estimate (e.g., a group of matches)
2. Compute transformation from seed group
3. Find *inliers* to this transformation
4. If the number of inliers is sufficiently large, re-compute least-squares estimate of transformation on all of the inliers

Keep the transformation with the largest number of inliers


RANSAC Line Fitting Example

- Task: Estimate the best line
 - *How many points do we need to estimate the line?*




RANSAC Line Fitting Example

- Task: Estimate the best line




RANSAC Line Fitting Example

- Task: Estimate the best line




RANSAC Line Fitting Example

- Task: Estimate the best line




RANSAC Line Fitting Example

- Task: Estimate the best line




RANSAC Line Fitting Example

- Task: Estimate the best line



RANSAC Line Fitting Example

- Task: Estimate the best line



Algorithm 15.4: RANSAC: fitting lines using random sample consensus

Determine:

- n — the smallest number of points required
- k — the number of iterations required
- t — the threshold used to identify a point that fits well
- d — the number of nearby points required
to assert a model fits well

Until k iterations have occurred

Draw a sample of n points from the data
uniformly and at random

Fit to that set of n points

For each data point outside the sample

Test the distance from the point to the line
against t ; if the distance from the point to the line
is less than t , the point is close

end

If there are d or more points close to the line
then there is a good fit. Refit the line using all
these points.

end

Use the best fit from this collection, using the
fitting error as a criterion

RANSAC: How many samples?



- How many samples are needed?
 - p is the probability of success
 - Suppose w is fraction of inliers (points from line).
 - n points needed to define hypothesis (2 for lines)
 - k samples chosen.
- Prob. that a single sample of n points is correct: w^n
- Prob. that all k samples fail is: $(1 - w^n)^k$
- Choose k high enough to keep this below desired failure rate. $k = \frac{\log(1-p)}{\log(1-w^n)}$

RANSAC: Computed k ($p=0.99$)

Sample size n	Proportion of outliers						
	5%	10%	20%	25%	30%	40%	50%
2	2	3	5	6	7	11	17
3	3	4	7	9	11	19	35
4	3	5	9	13	17	34	72
5	4	6	12	17	26	57	146
6	4	7	16	24	37	97	293
7	4	8	20	33	54	163	588
8	5	9	26	44	78	272	1177

After RANSAC

- RANSAC divides data into inliers and outliers and yields estimate computed from minimal set of inliers.
- Improve this initial estimate with estimation over all inliers (e.g. with standard least-squares minimization).
- But this may change inliers, so alternate fitting with re-classification as inlier/outlier.



RANSAC: Pros and Cons

- **Pros:**

- Robust against outliers
- A general method that can be applied to most cases
- Fast in huge data sets
- Easy to implement

- **Cons:**

- Has a certain probability of success
- Requires prior knowledge about the data
- Number of iterations increases logarithmically with outlier percentage

Summary

- **Fitting** problems require finding any supporting evidence for a model, even within clutter and missing features.
 - associate features with an explicit model
- If we know which points belong to the line, how do we find the “optimal” line parameters?
 - Least squares approaches
- What if there are outliers?
 - RANSAC
- What if there are many lines?
 - Voting methods: RANSAC, Hough transform
 - Voting approaches, such as the Hough transform, make it possible to find likely model parameters without searching all combinations of features.
 - Hough transform approaches for lines, circles, and other shapes

Next Time: Classification

Slide Credits

Some slides from Stanley Birchfield, Kristen Grauman, Svetlana Lazebnik, Jia-Bin Huang, Silvio Savarese, Steve Seitz, James Hays, Derek Hoiem, David Forsyth, Fei-Fei Li. Vivek Kwatra, Jinxiang Chai, David Lowe

Questions?